



Five-model embedding benchmark (70,425 posts, 256 tokens)


Coherence c_v




Coherence c_{npmi}




Topic diversity




Retrieval P@10



Retrieval MRR



Retrieval nDCG@10



MuRIL MentalBERT HingBERT-Mx-v2 HingRoBERTa-Mx BGE-m3